Meta Hackathon — Scope & Requirements

Oct 30, 2026 · @Mayar

What this is: Mayar is entering the Meta Global AI Developer Hackathon, a build competition for apps made with Meta's AI models (listed on the Cerebral Valley event page), with Mursal (Arabic for "sender"), a web app that helps small online sellers in the Gulf turn chat messages into clean orders. This document fixes what the project includes, what it deliberately leaves out and which rules it must obey, so the scope doesn't grow while it is being built. Terms are explained in the glossary at the end.

Goal

Mursal, an order assistant for sellers who take orders in Instagram and WhatsApp messages, reads a pasted customer chat, produces a structured order, drafts a polite reply asking for anything missing, and sends a payment invoice once the seller approves it.

- **Who it's for:** small sellers in the Gulf region who take orders in direct messages, where sizes, colours, addresses and payment details get lost in a long chat.
- **What winning means:** a working demo submitted to the Indie track (one of the hackathon's competition categories) by Wed 4 Nov 2026, 6:00 PM Riyadh time. Judges score Impact 30%, Use of Meta AI tooling 25%, Demo 25% and Technical depth 20%. The top 5 entries per track present live on Mon 9 Nov.
- **Separate from Altın:** Mayar is also entering the PayPal AI Hackathon with a different project called Altın. The two entries must not be the same project.
- **One test for every idea:** if a feature doesn't make the demo clearer or more reliable, it waits.

The story

A seller pastes a chat written in Gulf Arabic dialect, Modern Standard Arabic (the formal written Arabic), English or a mix, and gets back a clean order instead of scrolling a long thread to find the size, colour and address.

- **The seller in the demo:** a fictional persona, never a real person.
- **Hackathon theme:** "Build for Your World".

- **Product name:** Mursal (Arabic for "sender", pronounced MUR-sal). Chosen on 30 Oct after an independent review. Its trademark, app-store, domain and handle checks are still to do.
- **Evidence of impact:** if a real seller agrees, add one quote or 3 to 5 anonymised messages, with their permission. Otherwise the demo uses only chats labelled fictional.
- **After the hackathon:** a paid order inbox for Gulf social sellers, possibly as a subscription if Meta's platform rules allow. Whether a real market gap exists has not been researched yet (planned after 5 Nov).

What the first version includes

The demo needs six things, listed in build order. Everything runs in test mode, so no real money moves.

1. **Test chats and a live site:** at least 10 made-up customer chats (20 if time allows), each with the correct order written out by hand next to it, and the app hosted online on Render (a hosting service) as a website, an API and a database.
2. **Order extraction:** the server sends a pasted chat to Muse Spark (Meta's AI model, used through Meta's Model API). The reply must match a fixed order format, checked with zod (a TypeScript validation library). If it doesn't match, the server retries once and passes along the error. It returns the order plus a list of missing details.
3. **Ask-back reply:** when details are missing, the app drafts a short, polite message in the customer's dialect that asks only for what is missing. Nothing is sent automatically.
4. **Seller screens:** paste a chat, see an order card with the missing details flagged, edit any field, approve, and see a list of orders with their statuses. The screens work in Arabic (right to left) and English and are designed for phones first.
5. **Invoice on approval:** approving an order sends a PayPal sandbox invoice (PayPal's test mode) and the order shows when it has been paid. Payments sit behind a small interface so another provider could be added later.
6. **Accuracy test:** a Python script runs the labelled chats through Muse Spark and through one open-weights Meta model (a model whose weights are published) and reports accuracy, speed and cost per order, with one chart.

Stretch goal, dropped unless everything above is finished by the night of Mon 2 Nov: either SAM 3.1 (a Meta image model) cutting a product out of a seller's photo, or receiving one real message through a WhatsApp test number.

Left out on purpose

Anything below waits until after the submission. Mayar still needs to confirm this list.

- **Real Instagram or WhatsApp connections:** these need an app review by Meta, so the demo uses pasted chats. The stretch-goal test number is the only exception.
- **Sending anything to a customer automatically:** the seller always sends.
- **Real customer data:** fictional chats only.
- **Other payment providers:** the demo says that mada and Apple Pay (popular payment methods in the Gulf) come next, but only PayPal's test mode is built.
- **Real payments:** test mode only.
- **Accounts, logins and subscriptions:** the business case is for after the hackathon.
- **Anything from the Altın project:** its savings-circle features stay there.

Rules the product always follows

These hold no matter what the AI says. If a feature would break one, the feature loses.

1. **The seller stays in control.** The AI drafts; the seller edits and approves. Nothing reaches a customer, and no invoice is sent, without that approval.
2. **Model output is never trusted until it is checked.** Every reply is validated against the fixed order format, with one retry that includes the error.
3. **Fictional data only.** Every demo chat is labelled fictional, raw chats are not logged, and the README includes a short privacy note on what is stored and how to delete it.
4. **No secrets in the code or its history,** and the extraction endpoint is rate limited.
5. **The exact Muse Spark version is fixed** in the configuration and in the accuracy test, because newer versions keep appearing and would change the results.
6. **Payments sit behind a swappable interface,** and the product is a test-mode demo with no real payments.
7. **No Instagram or WhatsApp logos and no copyrighted music** in the demo video.

When it is done

The project is done when the demo runs start to finish on camera without a fix mid-take and the submission form is complete. The scene-by-scene script for the demo video is still to be written; the boxes below are what it must show.

- [ ] A messy Gulf-dialect chat goes in and an order card comes out, with the missing details flagged
- [ ] The ask-back reply is drafted in the customer's dialect
- [ ] The seller edits and approves, a PayPal sandbox invoice is sent, and the order flips to paid
- [ ] A chart shows Muse Spark against the open-weights model on the labelled chats
- [ ] The live demo link works
- [ ] The README covers the problem, the architecture, which Meta models are used and how, the accuracy table, steps for judges to test it, an MIT licence (a permissive open-source licence), and the line "sandbox demo, no real payments"
- [ ] The video is under 3 minutes (assumed; confirm in the Official Rules) with English narration or subtitles
- [ ] The submission picks the Indie track, links the code, video and live demo, names the Meta AI products used, and a screenshot of the confirmation is saved

Dates and constraints

Mayar submits by Wed 4 Nov at 6:00 PM Riyadh time, about eleven hours before the hard deadline, and spends about 90% of her working hours on this project from Fri 30 Oct to Wed 4 Nov.

What	When
Read the Official Rules and claim $150 of free Meta Model API credits	Before building; the credits form closes Thu 5 Nov, 5:00 AM Riyadh
Stretch goal is dropped unless everything else is finished	Mon 2 Nov, night
Checkpoint 1: a deployed endpoint returns a valid order from a real Muse Spark call	Sun 1 Nov, 9 to 10 PM
Submission target	Wed 4 Nov, 6:00 PM
Checkpoint 2: has this project been submitted?	Wed 4 Nov, 11:00 PM
Hard deadline	Thu 5 Nov, 5:00 AM Riyadh (Wed 4 Nov, 9:00 PM US Eastern)

What	When
Judging	4 to 8 Nov
Top 5 per track present live	Mon 9 Nov

- **Technology:** React and TypeScript for the website, a Node and TypeScript API, a Postgres database, all hosted on Render, plus a shared zod format definition and Python for the accuracy test.
- **Eligibility:** Saudi Arabia is not on the excluded list, but this must be confirmed in the Official Rules, including the clause about government-affiliated people. Prizes likely require a W-8BEN (a US tax form that non-US winners return), and up to about 30% US tax withholding is possible.
- **Unknown until the rules are read:** the submission format, the video length and the rule on reusing existing code, because the Official Rules can only be read after signing in.

Changing the scope

A new idea goes on a parking-lot list first and only joins the first version if it passes all three tests below. This rule is a proposal for Mayar to confirm.

1. It is needed to keep one of the rules above true, or to make a must-have demo step work.
2. Something of equal size leaves the first version to make room.
3. It happens before Checkpoint 1 on Sun 1 Nov.

Cuts agreed in advance (30 Oct):

- If Checkpoint 1 fails, the accuracy test shrinks to 10 chats and a results table that takes 45 to 60 minutes, the app stays on one screen, and all work on the PayPal entry moves to Thu 5 Nov.
- If this project isn't submitted by Checkpoint 2, the Altın entry drops Payouts (PayPal's feature for sending money to a person) and demos only invoices, payment notifications and reminders.
- If time runs short, the accuracy test is written in TypeScript first and ported to Python later.

Decisions still open for Mayar:

- [ ] Check Mursal against the USPTO and Saudi SAIP trademark databases, the app stores, domains and social handles
- [ ] Which open-weights Meta model the accuracy test compares against Muse Spark
- [ ] Stretch goal: SAM 3.1 cut-out, a WhatsApp test number, or neither
- [ ] Whether a real seller can supply a quote or a few messages, with permission
- [ ] Confirm or strike each item under "Left out on purpose"

Glossary

Term	Meaning
Muse Spark	Meta's AI language model, used here to read customer chats
Meta Model API	The service through which apps call Meta's AI models
Open-weights model	An AI model whose internal weights are published, so anyone can run it
SAM 3.1	A Meta model that cuts an object out of a photo
zod	A TypeScript library that checks data matches a fixed format
Render	A hosting service where the website, API and database run
PayPal sandbox	PayPal's test mode, where invoices and payments use fake money
Indie track	One of the hackathon's competition categories
Official Rules	The hackathon's legal rules, readable only after signing in
W-8BEN	A US tax form that non-US prize winners return before being paid
Altın	Mayar's separate entry in the PayPal AI Hackathon, an AI assistant for savings circles